## Jacobi's Method

Jacobi's method is an iterative, numerical method for solving a system of linear equations. Formally, given a full rank  $n \times n$  matrix  $A \in \mathbb{R}^{n \times n}$  and a vector  $b \in \mathbb{R}^n$ , Jacobi's method iteratively approximates  $x \in \mathbb{R}^n$  for:

$$Am = b$$

Given the matrix A,

$$A = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \cdots & a_{nn} \end{bmatrix}$$

we first separate A into its diagonal elements D and the remaining elements R such that A = D + R with:

$$D = \begin{bmatrix} a_{11} & 0 & \cdots & 0 \\ 0 & a_{22} & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & a_{700} \end{bmatrix}, \qquad R = \begin{bmatrix} 0 & a_{12} & \cdots & a_{1m} \\ a_{21} & 0 & \cdots & a_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & 0 \end{bmatrix}$$

Jacobi's method then follows the following steps till convergence or termination:

- 1. Initialize  $x: x \leftarrow \begin{bmatrix} 0 & 0 & \dots & 0 \end{bmatrix}$
- 2. D = diag(A)
- 3. R = A D
- 4. while ||Ax b|| > l do

(a) update 
$$x \leftarrow D^{-1}(b - Rx)$$

where ||x|| is the L2-norm, and l is a parameter for the termination accuracy.

A matrix A is diagonally dominant, if its diagonal elements are larger than the sum of absolutes of the corresponding rows, i.e., iff:

$$|a_{ij}| > \sum_{j \neq i} |a_{ij}|.$$

It can be shown that for diagonally dominant matrices, Jacobi's method is guaranteed to converge.

## Parallel Algorithm

Data distribution. We use a 2-dimensional grid as the communication network for this problem. We assume that the number of processors is a perfect square:  $p = q \times q$  with  $q = \sqrt{p}$ , arranged into a mesh of size  $q \times q$ .

The inputs are distributed on the grid as follows:

• The n-by-n matrix A is block distributed onto the grid. The rows of A are distributed onto the rows of the processor-grid such that either \(\bigcap\_q^2\) or \(\bigcap\_q^2\) rows of A are distributed onto each row of the grid. More specifically, the first \((n \text{ mod } q\)) rows of the grid contain \(\bigcap\_q^2\) rows of A, and the remaining rows of the grid contain \(\bigcap\_q^2\)] rows of A. The same applies to the distribution of columns. A processor with coordinates \((i, j)\) thus has the following size local matrix:

$$\begin{cases} \left\lceil \frac{nq}{q} \right\rceil^2 & \text{if } i < (n \mod q) \text{ and } j < (n \mod q) \\ \left\lceil \frac{nq}{q} \right\rceil \times \left\lceil \frac{nq}{q} \right\rceil & \text{if } i < (n \mod q) \text{ and } j \geq (n \mod q) \\ \left\lceil \frac{nq}{q} \right\rceil \times \left\lceil \frac{nq}{q} \right\rceil & \text{if } i \geq (n \mod q) \text{ and } j < (n \mod q) \\ \left\lceil \frac{nq}{q} \right\rceil^2 & \text{if } i \geq (n \mod q) \text{ and } j \geq (n \mod q) \end{cases}$$

The size n vector b and x are equally block distributed only along the first column of the
processor-grid, i.e., only among processors with indexes (i, 0). Processor (i, 0) will thus have
the following local size:

$$\begin{cases} \left\lceil \frac{n}{q} \right\rceil & \text{if } i < (n \mod q) \\ \left\lceil \frac{n}{q} \right\rceil & \text{if } i \ge (n \mod q) \end{cases}$$

Parallel Matrix-Vector Multiplication Let A be a n-by-n matrix and let x and y be n dimensional vectors. We wish to compute y = Ax. Assume that the square matrix A and the vector x are distributed among the processor grid as explained above. The final answer y will again be distributed in the same fashion as the vector x was. To do so, we will first "transpose" the vector x on the grid, such that a processor with index (i,j) will end up with the elements that were on processor (j,0) according to the above distribution. We do this by first sending the elements from a processor (i,0) to its corresponding diagonal processor (i,i), using a single MPI\_Send and the sub-communicator for the row of processors. We then broadcast the received elements from (i,i) along each column of the grid using a sub-communicator for the column of processors.

Now each processor has the elements it needs for its local matrix vector multiplication. We multiply the local vector with the local matrix and then use a parallel reduction to sum up the resulting vectors along the rows of the grid back onto the processors of the first column. The final result of the multiplication thus ends up distributed among the first column in the same way that the input x was.

Parallel Jacobi For parallelizing Jacobi's method, we distribute A and R as described above and use parallel matrix-vector multiplication for calculating Rx. The vectors x, and b, and the diagonal

elements D of A are distributed among the first column of the processor grid. We can thus update x by  $x_i \leftarrow \frac{1}{d_i} \langle b_i - \langle Rx \rangle_i \rangle$  using only local operations.

In order to detect termination, we calculate Ax using parallel matrix vector multiplication, and then subtract b locally on the first column of processors (where the result of the multiplication and b are located). Next, we calculate the L2-norm ||Ax - b|| by first calculating the sum of squares locally and then performing a parallel reduction along the first column of processors. We can now determine whether to terminate or not by checking the L2-norm against the termination criteria and then broadcasting the result along rows. Now every processor knows whether or not to continue with further iterations. In your implementation, you should not only check for the termination criteria, but also terminate after a pre-set maximum number of iterations.