

Parallel List Ranking

Advanced Algorithms & Data Structures

Lecture Theme 17

Prof. Dr. Th. Ottmann

Summer Semester 2006

Two parallel list ranking algorithms


- An $O(\log n)$ time and $O(n \log n)$ work list ranking algorithm.
- An $O(\log n \log \log n)$ time and $O(n)$ work list ranking algorithm.

List ranking

Input: A linked list L of n elements.

L is given in an array S such that the entry $S(i)$ contains the index of the node which is the successor of the node i in L .

Output: The distance of each node i from the end of the list.



Sequential list ranking algorithm

List ranking can be solved in $O(n)$ time sequentially for a list of length n .

Exercise : Design an algorithm.

- Hence, a work-optimal parallel algorithm should do only $O(n)$ work.

A simple parallel list ranking algorithm

Output: For each $1 \leq i \leq n$, the distance $R(i)$ of node i from the end of the list.

begin

for $1 \leq i \leq n$ do in parallel

if $S(i) \neq 0$ then $R(i) := 1$

else $R(i) := 0$

endfor

while $S(i) \neq 0$ and $S(S(i)) \neq 0$ do


Set $R(i) := R(i) + R(S(i))$

Set $S(i) := S(S(i))$

end


Initialization phase of the list ranking algorithm

- At the start of an iteration of the while loop, $R(i)$ counts the nodes in a sublist starting at i (a subset of nodes which are adjacent in the list).



Iteration phase

- After the iteration, $R(i)$ counts the nodes in a sublist of double the size.
- When the while loop terminates, $R(i)$ counts all the nodes starting from i and until the end of the list



Complexity and model

- The algorithm terminates after $O(\log n)$ iterations of the while loop.
- The work complexity is $O(n \log n)$ since we allocate one processor for each node.
- We need the CREW PRAM model since several nodes may try to read the same successor (S) values.

Exercise :

Modify the algorithm to run on the EREW PRAM with the same time and processor complexities.

Discussion of the algorithm

- Our aim is to modify the simple algorithm so that it does optimal $O(n)$ work.
- The best algorithm would be the one which does $O(n)$ work and takes $O(\log n)$ time.
- There is an algorithm meeting these criteria, however the algorithm and its analysis are very involved.
- We will study an algorithm which does $O(n)$ work and takes $O(\log n \log \log n)$ time.
- However, in future we will use the optimal algorithm for designing other algorithms.

The strategy for an optimal algorithm


1. Shrink the initial list L by removing some of the nodes.
The modified list should have $O(n / \log n)$ nodes.
2. Apply the pointer jumping technique (the suboptimal algorithm) on the list with $O(n / \log n)$ nodes.
3. Restore the original list and rank all the nodes removed in Step 1.

The important step is Step1. We need to choose a subset I of nodes for removal.

Independent sets

Definition

A set I of nodes is **independent** if whenever $i \in I$, $S(i) \notin I$.



The blue nodes form an independent set in this list

- The main task is to pick an independent set correctly.
- We pick an independent set by first coloring the nodes of the list by two colors.

2-coloring the nodes of a list

Definition: A k -coloring of a graph G is a mapping: $c: V \rightarrow \{0, 1, \dots, k - 1\}$ such that

$c(i) \neq c(j)$ if $\langle i, j \rangle \in E$.

- It is very easy to design an $O(n)$ time sequential algorithm for 2-coloring the nodes of a linked list.

We will assume the following result:

Theorem: A linked list with n nodes can be 2-colored in $O(\log n)$ time and $O(n)$ work.


Independent sets

- When we 2-color the nodes of a list, alternate nodes get the same color.
 - Hence, we can remove the nodes of the same color to reduce the size of the original list from n to $n/2$.
 - However, we need a list of size $\frac{n}{\log n}$ to run our pointer jumping algorithm for list ranking.
 - If we repeat the process $\log \log n$ time, we will reduce the size of the list to $\frac{n}{2^{\log \log n}}$
- i.e., to $\frac{n}{\log n}$

Preserving the information

- When we reduce the size of the list to $\frac{n}{\log n}$, we have lost a lot of information because the removed nodes are no longer present in the list.
- Hence, we have to put back the removed nodes in their original positions to correctly compute the ranks of all the nodes in the list.
- Note that we have removed the nodes in $O(\log \log n)$ iterations.
- So, we have to replace the nodes also in $O(\log \log n)$ iterations.

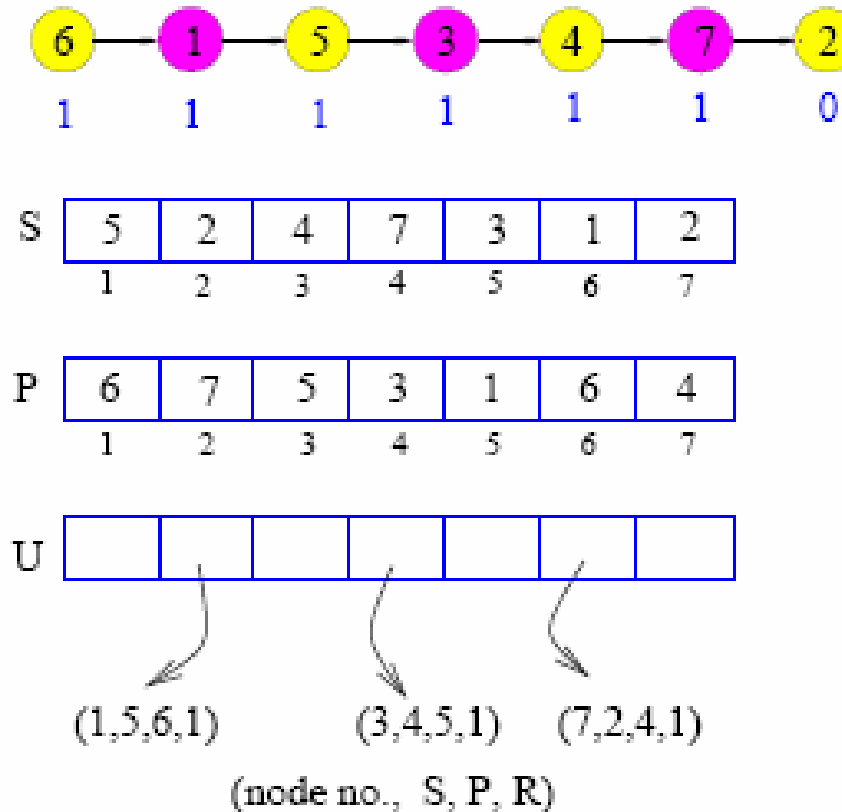
Preserving the information




- The **magenta** nodes are removed through **2**-coloring. The number of a node is shown in black.
- The initial rank of each node is shown below the node in the figure.
- The modified rank of an existing **yellow** node is shown in the second figure.
- The modified rank is with respect to the next existing node in the list.

Preserving the information

- P is the array which holds the predecessor indices.
- S is the array which holds the successor indices.
- i is the number of the removed node.
- U is the array where we store information for the removed nodes. This information is used later to restore the removed nodes.



Preserving the information



S

5	2	4	7	3	1	2
1	2	3	4	5	6	7

P

6	7	5	3	1	6	4
1	2	3	4	5	6	7

U

--	--	--	--	--	--	--

(1,5,6,1) (3,4,5,1) (7,2,4,1)

(node no., S, P, R)

Preserving the information

The code executed for removing the nodes in the independent set I :

for all $i \in I$ pardo


Set $U(i) := (i, S(i), P(i))$

Set $R(P(i)) := R(P(i)) + R(i)$

Set $S(P(i)) := S(i)$

Set $P(S(i)) := P(i)$

Preserving the information

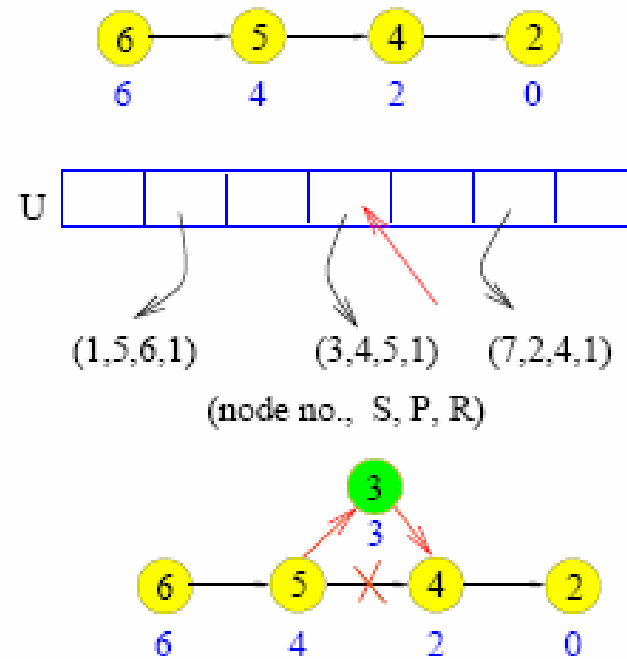


- The predecessor and successor pointers are updated
- The ranks of nodes which are immediately before a removed node are updated.

Preserving the information

- The array P and S can be compacted in $O(n)$ work and $O(\log n)$ time after each phase of removal of nodes. This is done through a parallel prefix computation.
- Recall that we have to do the removal of nodes for $\log \log n$ phases to reduce the number of nodes to $\frac{n}{\log n}$.
- We can keep a separate U array for each phase.

Putting back a node



- A removed node can be reinserted in $O(1)$ time

Complexity

- 2-coloring in each stage can be done in $O(\log n)$ time and $O(n)$ work.
- The compaction of the P and S arrays is done through a parallel prefix computation in $O(\log n)$ time and $O(n)$ work.
- There are in all $O(\log \log n)$ stages in which we have to remove nodes. Each stage takes $O(\log n)$ time due to the 2-coloring and parallel prefix computations.
- At the end, we do one list ranking for a list of size $O(n/\log n)$. This takes $O(n)$ work and $O(\log n)$ time.
- So, the overall complexity is $O(n)$ work and $O(\log n \log \log n)$ time.